

# Forecasting: principles and practice

Lab Session 2

23 September 2014

---

Before doing any exercises in R, load the **fpp** package using `library(fpp)`.

1. The function `tsdisplay(data, plot.type="scatter")` is useful for showing a time plot, ACF plot and lagged scatterplot on the same graph. Use it to produce plots of the following time series:

`bricksq, hsales, ibmclose`

Can you spot the effects of seasonality, cyclicity and trend?

2. For each of the same series (`bricksq`, `ibmclose`, `hsales`):
  - (a) Use either the naive or seasonal naive forecasting method and apply it to the full data set.
  - (b) Compute the residuals and plot their ACF. Do the residuals appear to be white noise? What did your forecasting method miss?
  - (c) Do a Ljung-Box test on the residuals. What do the results mean?

3. For the data set `bricksq`:

- (a) Split the data into two parts using

```
bricks1 <- window(bricksq, end=1987.99)
bricks2 <- window(bricksq, start=1988)
```

- (b) Check that your data have been split appropriately by producing the following plot.

```
plot(bricksq)
lines(bricks1,col="red")
lines(bricks2,col="blue")
```

- (c) Calculate forecasts using each of the four benchmark methods applied to `bricks1`.
- (d) Compare the accuracy of your forecasts against the actual values stored in `bricks2`. For example:

```
f1 <- meanf(bricks1)
accuracy(f1,bricks2)
```

- (e) Which method does best? Why?
- (f) For the best method, compute the residuals and plot them. For example

```
res <- residuals(f1)
plot(res)
hist(res, breaks="FD")
Acf(res)
```

Do the residuals appear to be uncorrelated and normally distributed?

4. Consider the daily closing IBM stock prices (data set `ibmclose`).
  - (a) Produce some plots of the data in order to become familiar with it.
  - (b) Split the data into a training set of 300 observations and a test set of 69 observations.
  - (c) Try various benchmark methods to forecast the training set and compare the results on the test set. Which method did best?
  - (d) For the best method, compute the residuals and plot them. What do the plots tell you?
  - (e) Can you invent a better forecasting method than any of the benchmark methods for these data?
5. Consider the sales of new one-family houses in the USA (Jan 1987 – Nov 1995). Data set: `hsales`.
  - (a) Produce some plots of the data in order to become familiar with it.
  - (b) Split the data into a training set and a test set, where the test set is the last two years of data.
  - (c) Try various benchmark methods to forecast the training set and compare the results on the test set. Which method did best?
  - (d) For the best method, compute the residuals and plot them. What do the plots tell you?
  - (e) Can you invent a better forecasting method than any of the benchmark methods for these data?